

# ML/DL for Everyone with **PYTORCH**

## Lecture 10: Basic CNN

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST

Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>

Videos: <http://bit.ly/PyTorchVideo>



# Call for Comments

Please feel free to add comments directly on these slides.

Other slides: <http://bit.ly/PyTorchZeroAll>



# ML/DL for Everyone with **PYTORCH**

## Lecture 10: Basic CNN

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST


Code: <https://github.com/hunkim/PyTorchZeroToAll>


Slides: <http://bit.ly/PyTorchZeroAll>

Videos: <http://bit.ly/PyTorchVideo>




# CNN







# Convolution in Action






# Simple convolution layer

Stride: 1x1



# Simple convolution layer




Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, No Padding








# Simple convolution layer

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, No Padding

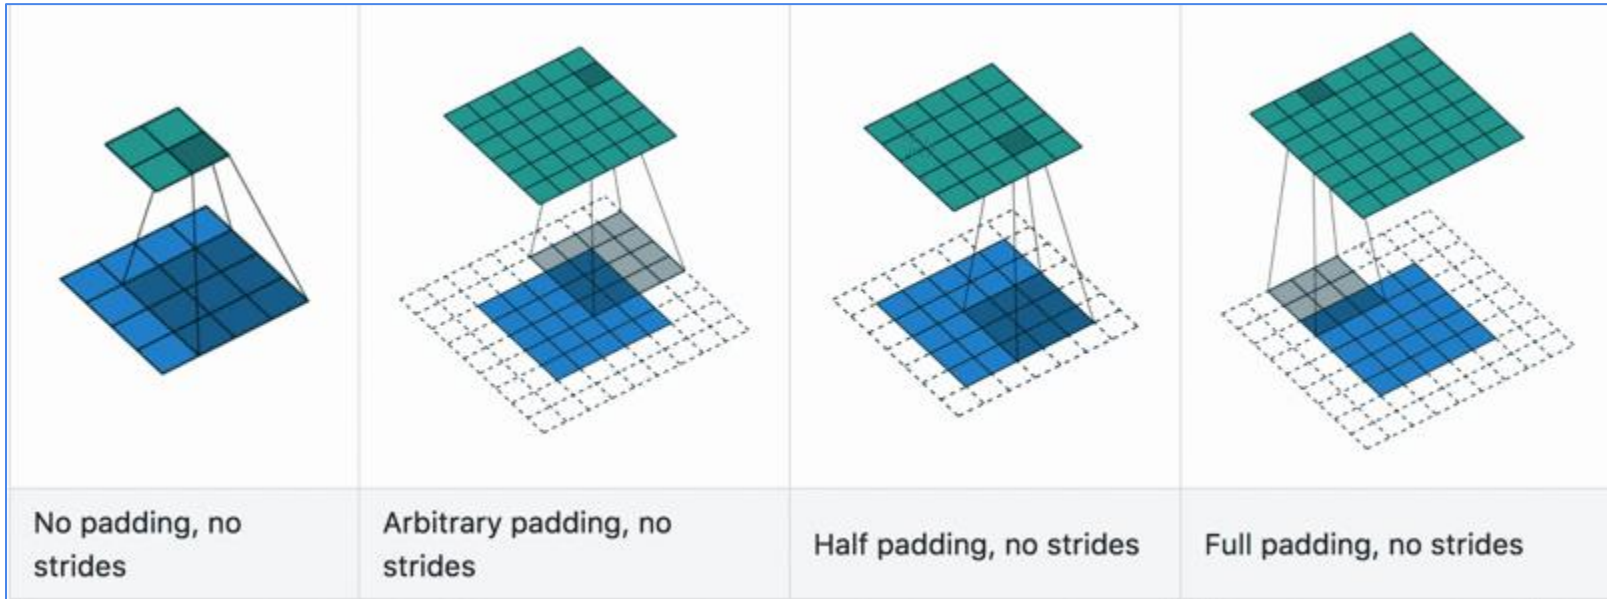


# Simple convolution layer


Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, With padding



# Convolution with padding in Action




# Convolution with stride in Action



# Convolution Layer

32x32x3 image



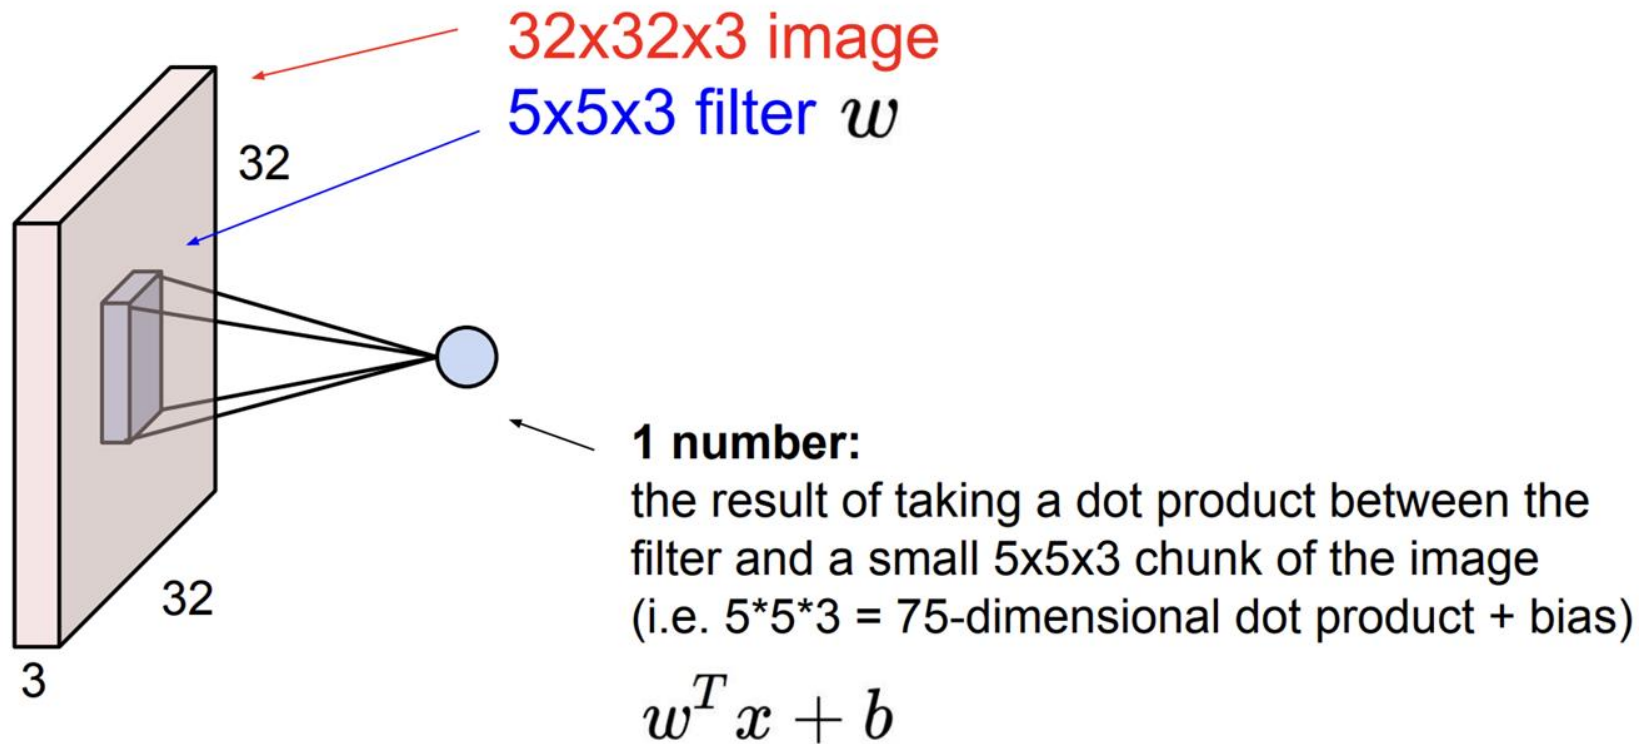
Filters always extend the full depth of the input volume

5x5x3 filter




**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

# Convolution Layer

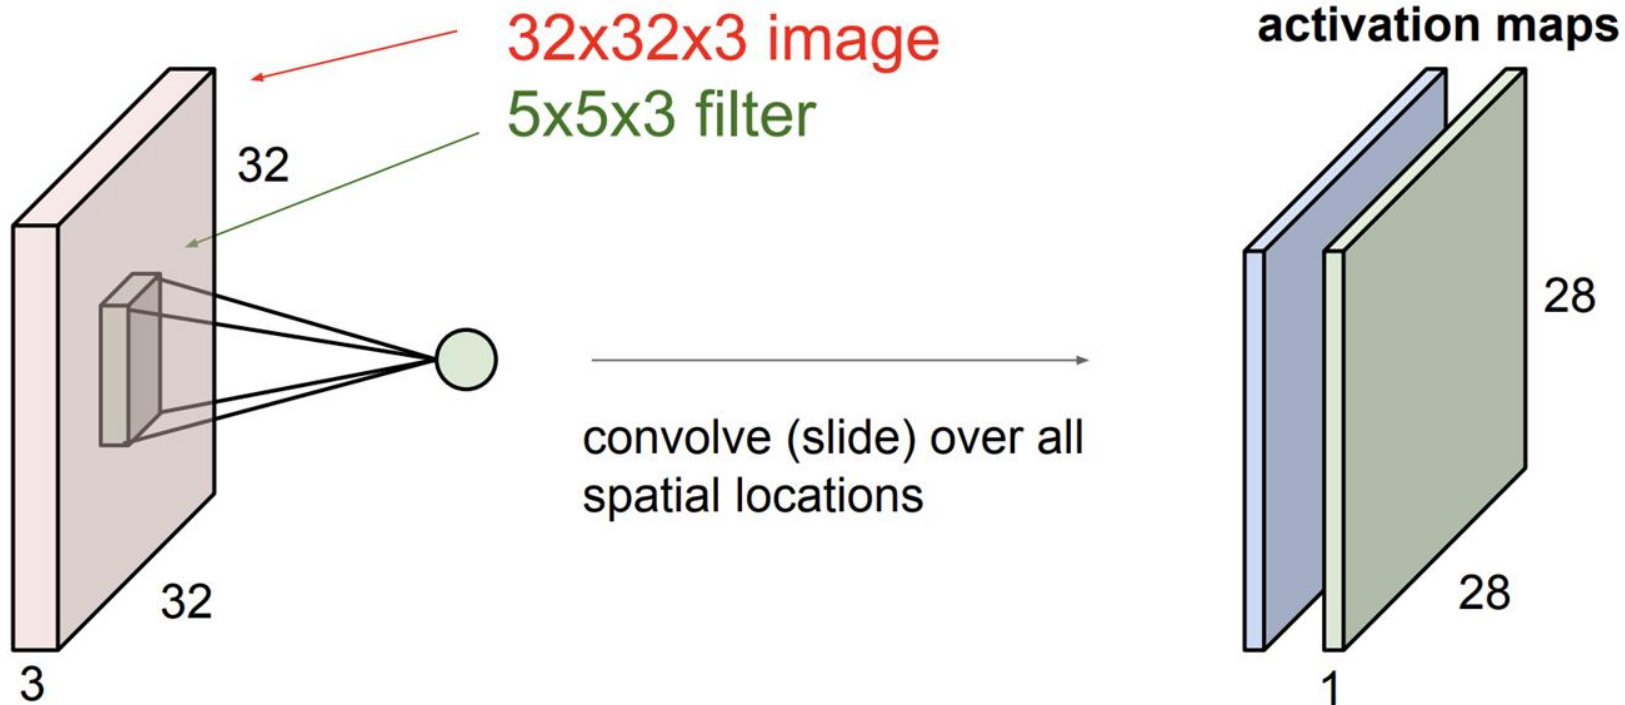


# Convolution Layer




# Convolution Layer

consider a second, **green** filter






For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:




We stack these up to get a “new image” of size 28x28x6!


**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



**Preview:** ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



# Max pooling



Single depth slice

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters  
and stride 2



6	8
3	4

# Max Pooling in Action

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1


3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

# Avg Pooling in Action

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1


1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3

# CNN




# Locally Connected Features

## FULLY CONNECTED NEURAL NET




## LOCALLY CONNECTED NEURAL NET

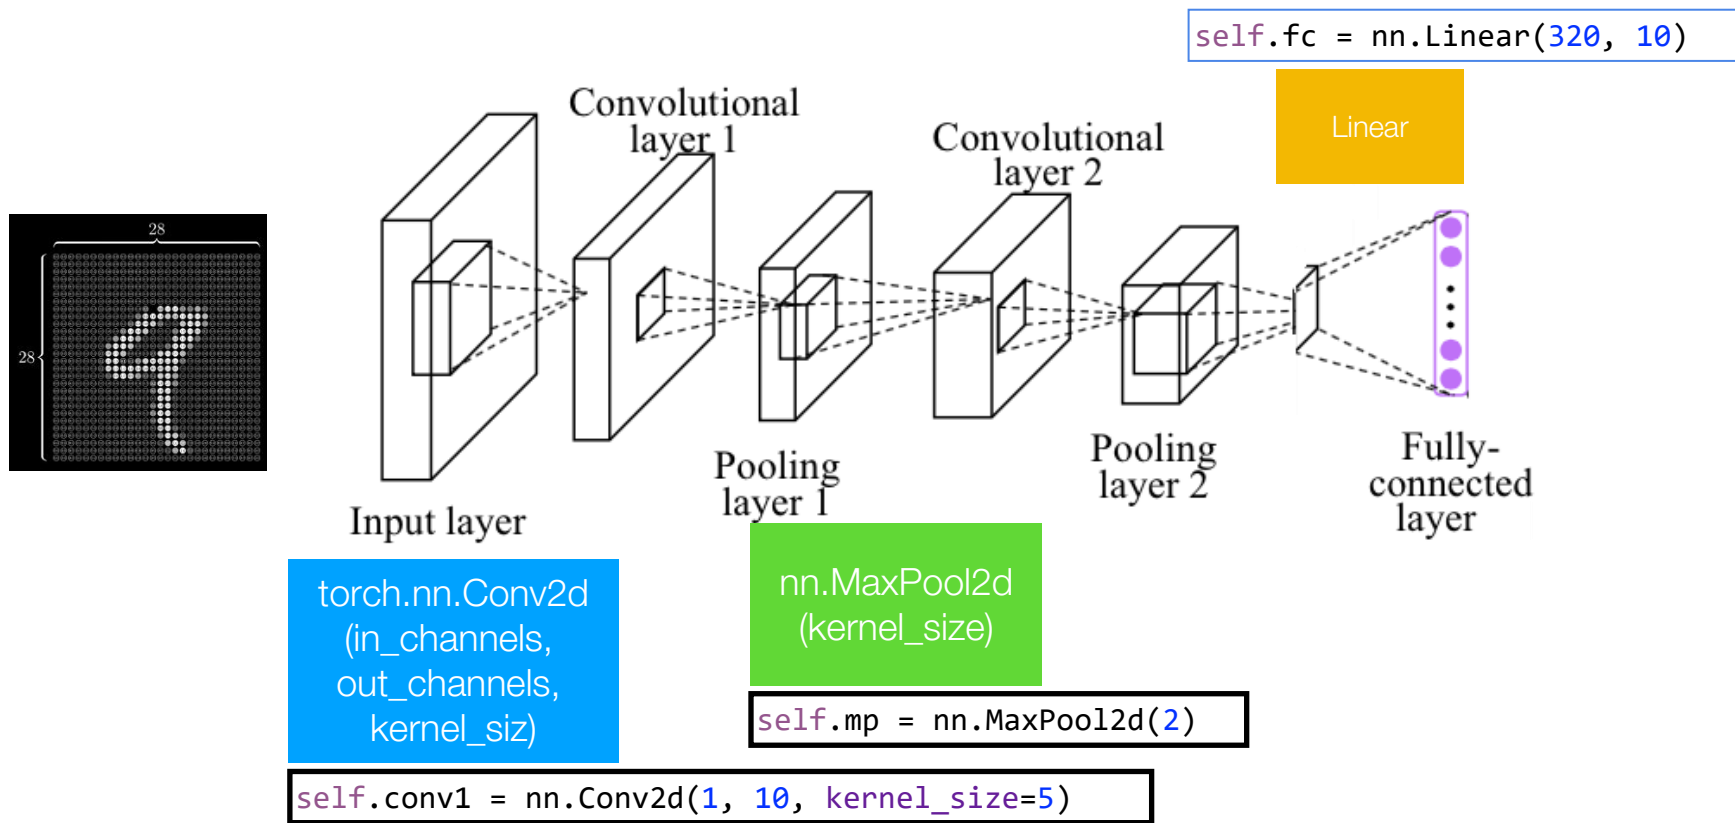





# Simple CNN



# Simple CNN





# Simple CNN




```
class Net(nn.Module):
```

```

def __init__(self):
    super(Net, self).__init__()
    self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
    self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
    self.mp = nn.MaxPool2d(2)
    self.fc = nn.Linear(100???, 10) # ??? -> 10

def forward(self, x):
    in_size = x.size(0)
    x = F.relu(self.mp(self.conv1(x)))
    x = F.relu(self.mp(self.conv2(x)))
    x = x.view(in_size, -1) # flatten the tensor
    x = self.fc(x)
    return F.log_softmax(x)

```



# Simple CNN




```
class Net(nn.Module):
```

```
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(100???, 10) # ??? -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

```
RuntimeError: size mismatch, m1: [64 x 320], m2: [100 x 10]
```



# Simple CNN




```
class Net(nn.Module):
```

```
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10) # 320 -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

```
RuntimeError: size mismatch, m1: [64 x 320], m2: [100 x 10]
```



# Simple CNN



```
class Net(nn.Module):
```


```
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10) # 320 -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

Train Epoch: 9 [46080/60000 (77%)]	Loss: 0.108415
Train Epoch: 9 [46720/60000 (78%)]	Loss: 0.140700
Train Epoch: 9 [47360/60000 (79%)]	Loss: 0.090830
Train Epoch: 9 [48000/60000 (80%)]	Loss: 0.031640
Train Epoch: 9 [48640/60000 (81%)]	Loss: 0.014934
Train Epoch: 9 [49280/60000 (82%)]	Loss: 0.090210
Train Epoch: 9 [49920/60000 (83%)]	Loss: 0.074975
Train Epoch: 9 [50560/60000 (84%)]	Loss: 0.058671
Train Epoch: 9 [51200/60000 (85%)]	Loss: 0.023464
Train Epoch: 9 [51840/60000 (86%)]	Loss: 0.018025
Train Epoch: 9 [52480/60000 (87%)]	Loss: 0.098865
Train Epoch: 9 [53120/60000 (88%)]	Loss: 0.013985
Train Epoch: 9 [53760/60000 (90%)]	Loss: 0.070476
Train Epoch: 9 [54400/60000 (91%)]	Loss: 0.065411
Train Epoch: 9 [55040/60000 (92%)]	Loss: 0.028783
Train Epoch: 9 [55680/60000 (93%)]	Loss: 0.008333
Train Epoch: 9 [56320/60000 (94%)]	Loss: 0.020412
Train Epoch: 9 [56960/60000 (95%)]	Loss: 0.036749
Train Epoch: 9 [57600/60000 (96%)]	Loss: 0.163087
Train Epoch: 9 [58240/60000 (97%)]	Loss: 0.117539
Train Epoch: 9 [58880/60000 (98%)]	Loss: 0.032256
Train Epoch: 9 [59520/60000 (99%)]	Loss: 0.026360

Test set: Average loss: 0.0483, Accuracy: 9846/10000 (98%)

# Exercise 10-1: Implement CNN more layers



**WHAT  
NEXT?**



## **Lecture 11: Advanced CNN**