

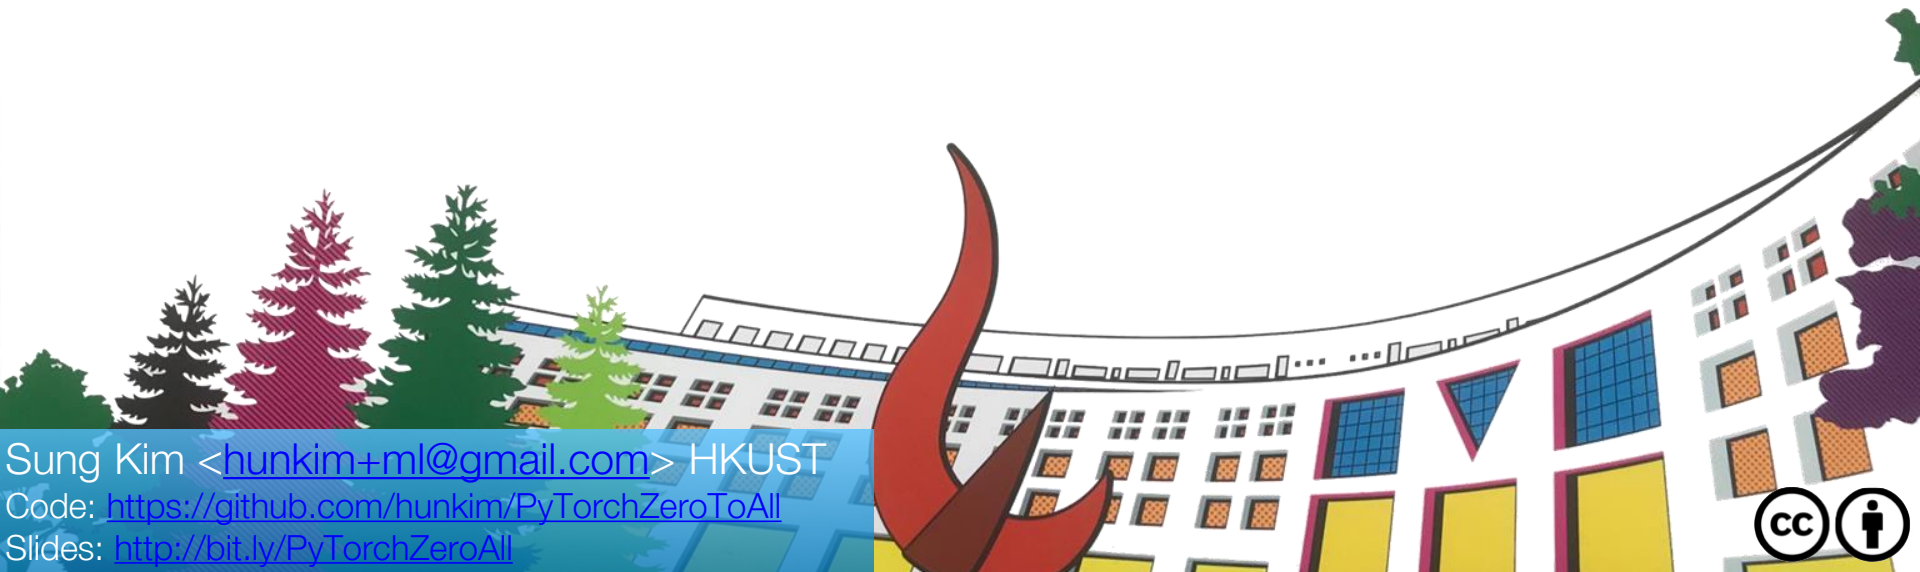
ML/DL for Everyone with **PYTORCH**

Lecture 14: Sequence to Sequence

Sung Kim <hunkim+ml@gmail.com> HKUST

Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>




Call for Comments

Please feel free to add comments directly on these slides.


Other slides: <http://bit.ly/PyTorchZeroAll>



DNN, CNN, RNN




RNN




CrossEntropy

RNN Applications


one to one




one to many




many to one




many to many



many to many



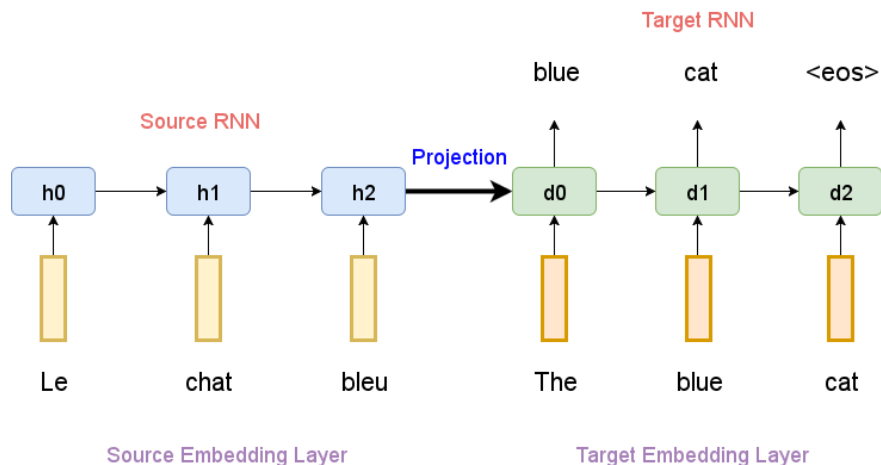
Sequence to Sequence




Lecture (TBA)

Implement Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation:

<https://arxiv.org/abs/1406.1078>



Summary: RNN (WIP)



Algorithm 1 RNN training loop

```
1: procedure TRAINING LOOP
2:    $h_0 = 0$ 
3:    $in_0 = 0$ 
4:    $l = 0$ 
5:   for each  $t$  from 1 to  $T$  do
6:      $h_t, out_t = \text{RNN-Module}(h_{t-1}, in_{t-1})$ 
7:      $l = l + \text{loss}(out_t, y_t)$ 
8:      $in_t = out_t$ 
9:   end for
10: end procedure
```

Modeling: $p(y)$

$$\log p(y) = \sum \log p(y_t | y_{<t})$$

$$h_t = f(h_{t-1}, y_{t-1})$$


$$p(y_t | y_{<t}) = \text{softmax}(g(h_t))$$

```
loss = 0
hidden = model.init_hidden()
Input = one_hot(labels[0])

for label in labels:
    hidden, output = model(hidden, input)
    loss += criterion(output, label)
    input = one_hot(output.max(1))

loss.backward()
optimizer.step()
```


Summary: S2S (WIP)



Algorithm 2 Sequence to Sequence training loop

```
1: procedure TRAINING LOOP
2:    $h^e = \text{encoder}(x)$ 
3:    $h_0 = h_l^e$ 
4:    $in_0 = \text{SOS}$ 
5:    $l = 0$ 
6:   for each  $t$  from 1 to  $T$  do
7:      $h_t, out_t = \text{decoder}(h_{t-1}, in_{t-1})$ 
8:      $l = l + \text{loss}(out_t, y_t)$ 
9:      $in_t = out_t$ 
10:  end for
11: end procedure
```

Modeling: $p(y|x)$

$h^e = \text{encoder}(x)$

$$\log p(y|x) = \sum \log p(y_t | y_{<t}, h^e)$$

$$h_t = f(h_{t-1}, y_{t-1}, h^e)$$

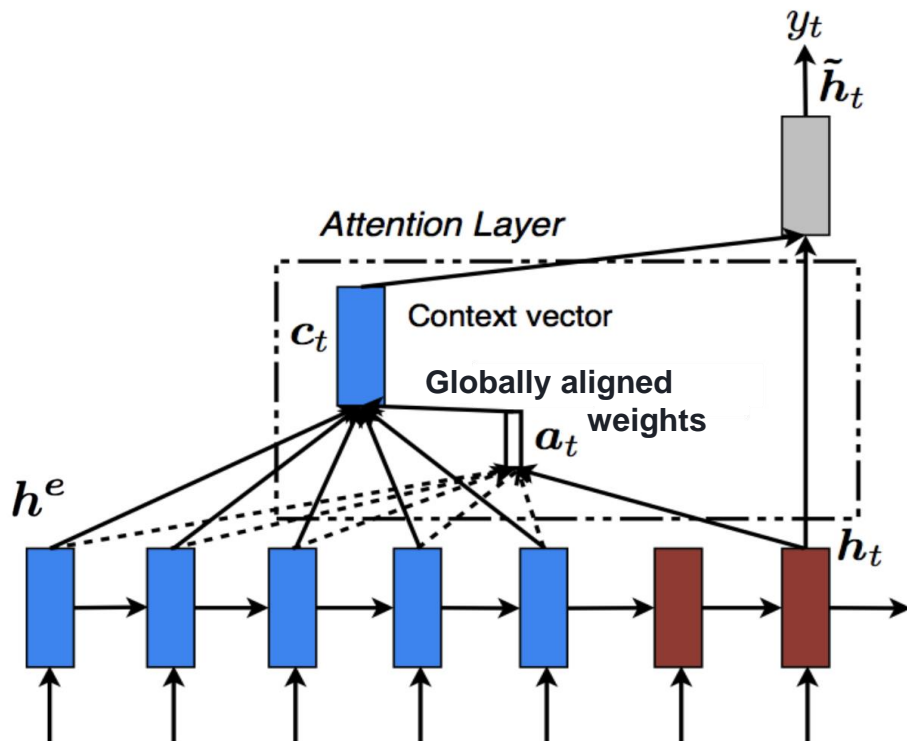
$$p(y_t | y_{<t}, h^e) = \text{softmax}(g(h_t))$$

```
loss = 0
hidden = encoder(x)
input = SOS

for label in labels:
    hidden, output = model(hidden, input)
    loss += criterion(output, label)
    input = one_hot(output.max(1))

loss.backward()
optimizer.step()
```

Summary: Attention (WIP)



Modeling $p(y|x)$

$h^e = \text{encoder}(x)$

$$\log p(y|x) = \sum \log p(y_t|y_{<t}, h^e)$$

$$h_t = f(h_{t-1}, y_{t-1}, h^e)$$

$$\tilde{h}_t = \tanh(W_c[c_t; h_t])$$

$$p(y_t|y_{<t}, h^e) = \text{softmax}(g(\tilde{h}_t))$$

$$c_t = h^e a_t$$

$$a_t = \text{align}(h_t, h^e) = \frac{\exp(\text{score}(h_t, h^e))}{\sum_i \exp(\text{score}(h_t, h_i^e))}$$

$$\text{score}(h_t, h^e) = h_t^T W_a h^e$$

Algorithm 3 Attention training loop

```

1: procedure TRAINING LOOP
2:    $h^e = \text{encoder}(x)$ 
3:    $h_0 = h_l^e$ 
4:    $in_0 = \text{SOS}$ 
5:    $l = 0$ 
6:   for each  $t$  from 1 to  $T$  do
7:      $s_t = h_t^T W_a h^e$ 
8:      $a_t = \text{softmax}(s_t)$ 
9:      $c_t = h^e a_t$ 
10:     $h_t, out_t = \text{decoder}(h_{t-1}, in_{t-1}, c_t)$ 
11:     $l = l + \text{loss}(out_t, y_t)$ 
12:     $in_t = out_t$ 
13:   end for
14: end procedure

```

References

- Sequence to Sequence
 - Sequence to Sequence models:
<https://github.com/MaximumEntropy/Seq2Seq-PyTorch>
 - Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation: <https://arxiv.org/abs/1406.1078>
- Attention Models
 - Attention and Augmented Recurrent Neural Networks
<https://distill.pub/2016/augmented-rnns/>
 - Neural Machine Translation by Jointly Learning to Align and Translate:
<https://arxiv.org/abs/1409.0473>
 - Effective Approaches to Attention-based Neural Machine Translation:
<https://arxiv.org/abs/1508.04025>

Exercise 13-1

Implement Neural Machine Translation by Jointly Learning to Align and Translate: <https://arxiv.org/abs/1409.0473>




Figure 1: The graphical illustration of the proposed model trying to generate the t -th target word y_t given a source sentence (x_1, x_2, \dots, x_T) .

$$s_i = f(s_{i-1}, y_{i-1}, c_i).$$

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j.$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})},$$

$$e_{ij} = a(s_{i-1}, h_j)$$


alignment model a as a feedforward neural network

<https://arxiv.org/abs/1409.0473>

Exercise 13-2

Implement A Neural Image Caption Generator:

<https://arxiv.org/abs/1411.4555>



If you've got this far, you did Good job! Congratulations!!

Interested in DL/ML related PHD, Postdoc at HKUST
and/or internship, residency, research fellows at LINE/NAVER?
Please email your exercises (Lectures 10 to 13) and CV to
hunkim+jobs@gmail.com.





Lecture 14: NSML, Smartest ML Platform