

Chapter 11 – Virtual Memory Management

Outline

- 11.1 Introduction
- 11.2 Locality
- 11.3 Demand Paging
- 11.4 Anticipatory Paging
- 11.5 Page Replacement
- 11.6 Page Replacement Strategies
 - 11.6.1 Random Page Replacement
 - 11.6.2 First-In-First-Out (FIFO) Page Replacement
 - 11.6.3 FIFO Anomaly
 - 11.6.4 Least-Recently-Used (LRU) Page Replacement
 - 11.6.5 Least-Frequently-Used (LFU) Page Replacement
 - 11.6.6 Not-Used-Recently (NUR) Page Replacement
 - 11.6.7 Modifications to FIFO: Second-Chance and Clock Page Replacement
 - 11.6.8 Far Page Replacement
- 11.7 Page-Fault-Frequency (PFF) Page Replacement



Chapter 11 – Virtual Memory Management

Outline (continued)

11.8 Page Release

11.9 Page Size

11.10 Program Behavior under Paging



Objectives

- After reading this chapter, you should understand:
 - the benefits and drawbacks of demand and anticipatory paging.
 - the challenges of page replacement.
 - several popular page-replacement strategies and how they compare to optimal page replacement.
 - the impact of page size on virtual memory performance.
 - program behavior under paging.



11.1 Introduction

- Replacement strategy
 - Technique a system employs to select pages for replacement when memory is full
 - Determines where in main memory to place an incoming page or segment
- Fetch strategy
 - Determines when pages or segments should be loaded into main memory
 - Anticipatory fetch strategies
 - Use heuristics to predict which pages a process will soon reference and load those pages or segments



11.2 Locality

- Process tends to reference memory in highly localized patterns
 - In paging systems, processes tend to favor certain subsets of their pages, and these pages tend to be adjacent to one another in process's virtual address space




11.3 Demand Paging

- Demand paging
 - When a process first executes, the system loads into main memory the page that contains its first instruction
 - After that, the system loads a page from secondary storage to main memory only when the process explicitly references that page
 - Requires a process to accumulate pages one at a time



11.3 Demand Paging

Figure 11.1 Space-time product under demand paging.



11.4 Anticipatory Paging

- Anticipatory paging
 - Operating system attempts to predict the pages a process will need and preloads these pages when memory space is available
 - Anticipatory paging strategies
 - Must be carefully designed so that overhead incurred by the strategy does not reduce system performance



11.5 Page Replacement

- When a process generates a page fault, the memory manager must locate referenced page in secondary storage, load it into page frame in main memory and update corresponding page table entry
- Modified (dirty) bit
 - Set to 1 if page has been modified; 0 otherwise
 - Help systems quickly determine which pages have been modified
- Optimal page replacement strategy (OPT or MIN)
 - Obtains optimal performance, replaces the page that will not be referenced again until furthest into the future



10.6 Page-Replacement Strategies

- A page-replacement strategy is characterized by
 - Heuristic it uses to select a page for replacement
 - The overhead it incurs



10.6.1 Random Page Replacement

- Random page replacement
 - Low-overhead page-replacement strategy that does not discriminate against particular processes
 - Each page in main memory has an equal likelihood of being selected for replacement
 - Could easily select as the next page to replace the page that will be referenced next



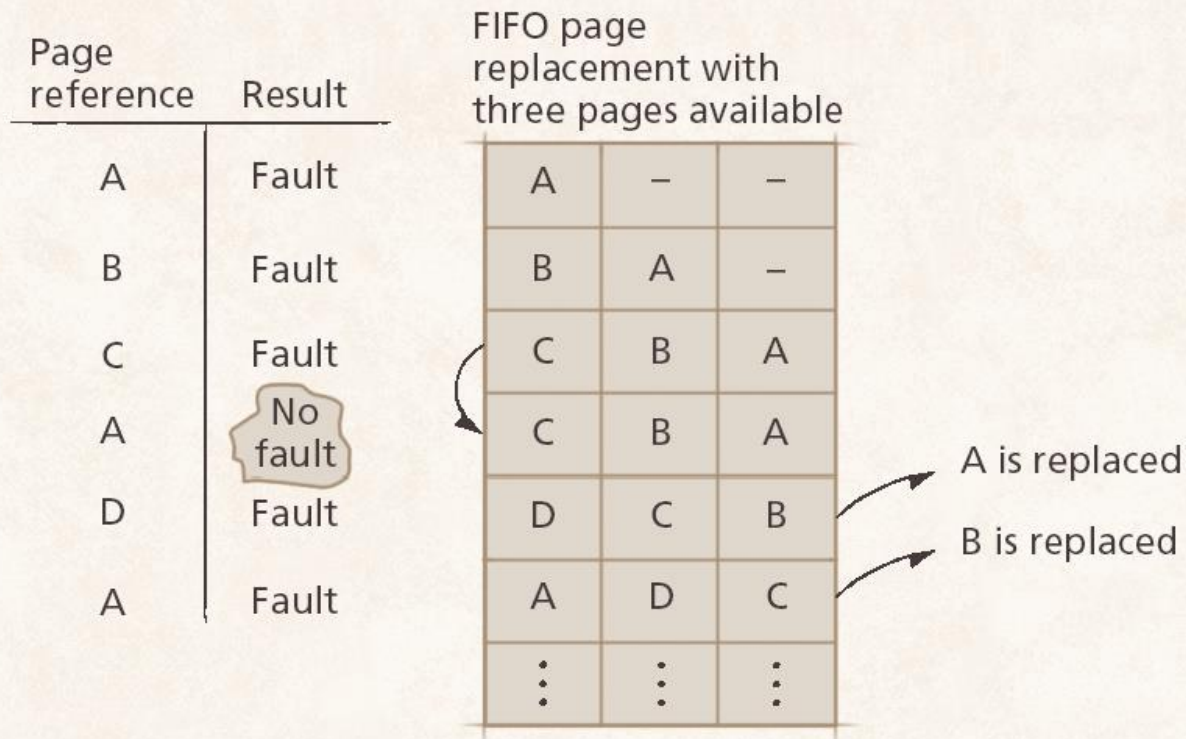
10.6.2 First-In-First-Out (FIFO) Page Replacement

- FIFO page replacement
 - Replace page that has been in the system the longest
 - Likely to replace heavily used pages
 - Can be implemented with relatively low overhead
 - Impractical for most systems



11.6.2 First-In-First-Out (FIFO) Page Replacement

Figure 11.2 First-in-first-out (FIFO) page replacement.



11.6.3 FIFO Anomaly

- Belady's (or FIFO) Anomaly
 - Certain page reference patterns actually cause more page faults when number of page frames allocated to a process is increased



11.6.3 FIFO Anomaly

Figure 11.3 FIFO anomaly—page faults can increase with page frame allocation.

Page reference	Result	FIFO page replacement with three pages available				FIFO page replacement with four pages available			
A	Fault	A	—	—	Fault	A	—	—	—
B	Fault	B	A	—	Fault	B	A	—	—
C	Fault	C	B	A	Fault	C	B	A	—
D	Fault	D	C	B	Fault	D	C	B	A
A	Fault	A	D	C	No fault	D	C	B	A
B	Fault	B	A	D	No fault	D	C	B	A
E	Fault	E	B	A	Fault	E	D	C	B
A	No fault	E	B	A	Fault	A	E	D	C
B	No fault	E	B	A	Fault	B	A	E	D
C	Fault	C	E	B	Fault	C	B	A	E
D	Fault	D	C	E	Fault	D	C	B	A
E	No fault	D	C	E	Fault	E	D	C	B
		Three “no faults”			Two “no faults”				



11.6.4 Least-Recently-Used (LRU) Page Replacement

- LRU page replacement
 - Exploits temporal locality by replacing the page that has spent the longest time in memory without being referenced
 - Can provide better performance than FIFO
 - Increased system overhead
 - LRU can perform poorly if the least-recently used page is the next page to be referenced by a program that is iterating inside a loop that references several pages



11.6.4 Least-Recently-Used (LRU) Page Replacement

Figure 11.4 Least-recently-used (LRU) page-replacement strategy.

Page reference	Result	LRU page replacement with three pages available		
A	Fault	A	–	–
B	Fault	B	A	–
C	Fault	C	B	A
B	No fault	B	C	A
B	No fault	B	C	A
A	No fault	A	B	C
D	Fault	D	A	B
A	No fault	A	D	B
B	No fault	B	A	D
F	Fault	F	B	A
B	No fault	B	F	A



11.6.5 Least-Frequently-Used (LFU) Page Replacement

- LFU page replacement
 - Replaces page that is least intensively referenced
 - Based on the heuristic that a page not referenced often is not likely to be referenced in the future
 - Could easily select wrong page for replacement
 - A page that was referenced heavily in the past may never be referenced again, but will stay in memory while newer, active pages are replaced



11.6.6 Not-Used-Recently (NUR) Page Replacement

- NUR page replacement
 - Approximates LRU with little overhead by using referenced bit and modified bit to determine which page has not been used recently and can be replaced quickly
 - Can be implemented on machines that lack hardware referenced bit and/or modified bit



11.6.6 Not-Used-Recently (NUR) Page Replacement

Figure 11.5 Page types under NUR.

<i>Group</i>	<i>Referenced</i>	<i>Modified</i>	<i>Description</i>
Group 1	0	0	Best choice to replace
Group 2	0	1	[Seems unrealistic]
Group 3	1	0	
Group 4	1	1	Worst choice to replace



11.6.7 Modification to FIFO: Second-Chance and Clock Page Replacement

- Second chance page replacement
 - Examines referenced bit of the oldest page
 - If it's off
 - The strategy selects that page for replacement
 - If it's on
 - The strategy turns off the bit and moves the page to tail of FIFO queue
 - Ensures that active pages are the least likely to be replaced
- Clock page replacement
 - Similar to second chance, but arranges the pages in circular list instead of linear list



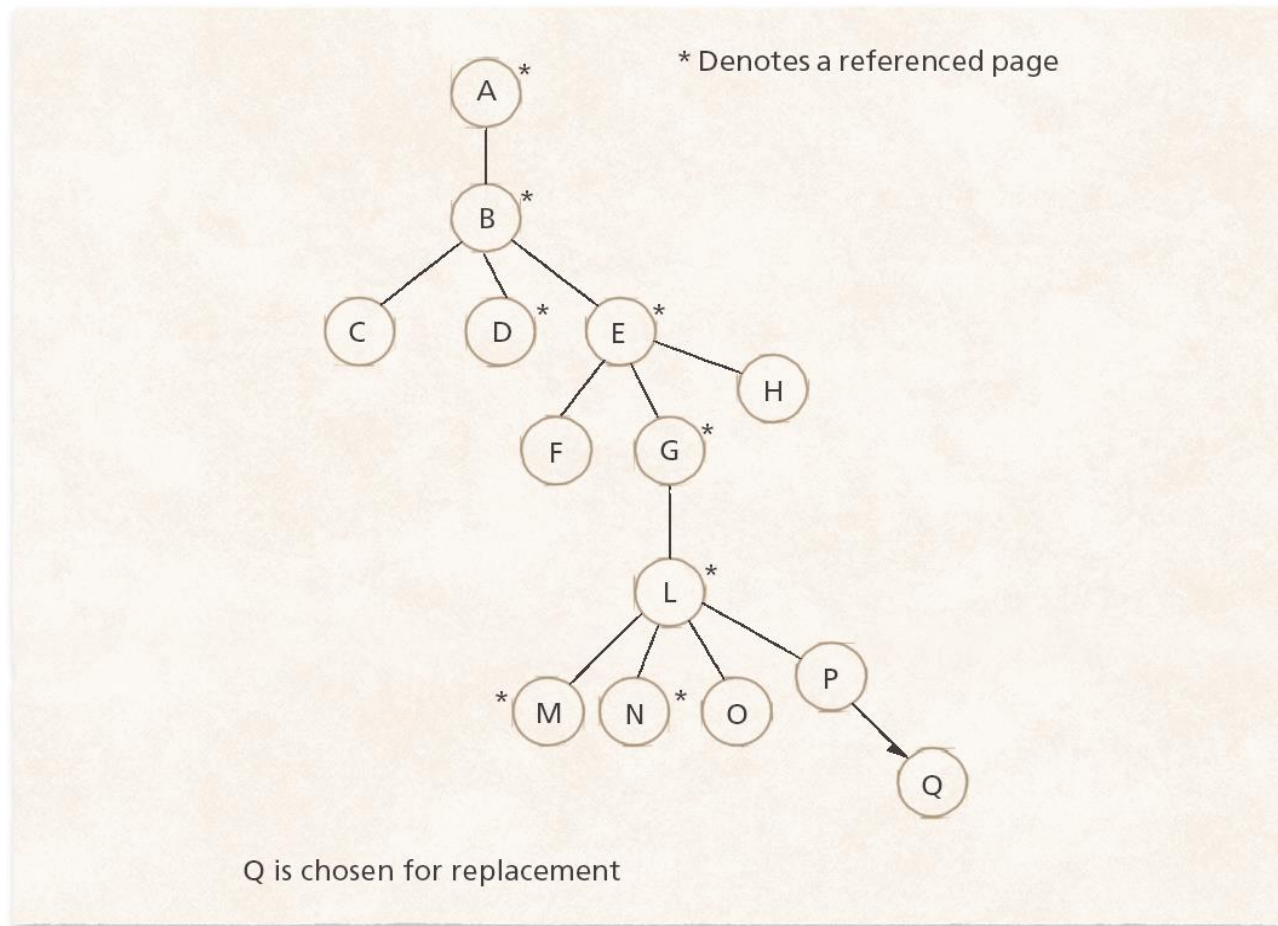
11.6.8 Far Page Replacement

- Far page replacment
 - Creates an access graph that characterizes a process's reference patterns
 - Replace the unreferenced page that is furthest away from any referenced page in the access graph
 - Performs at near-optimal levels
 - Has not been implemented in real systems
 - Access graph is complex to search and manage without hardware support



11.6.8 For Page Replacement

Figure 11.6 Far page-replacement-strategy access graph.



11.7 Page-Fault-Frequency (PFF) Page Replacement

- Adjusts a process's resident page set
 - Based on frequency at which the process is faulting
 - Based on time between page faults, called the process's interfault time
- Advantage of PFF over working set page replacement
 - Lower overhead
 - PFF adjusts resident page set only after each page fault
 - Working set mechanism must operate after each memory reference



11.8 Page Release

- Inactive pages can remain in main memory for a long time until the management strategy detects that the process no longer needs them
 - One way to solve the problem
 - Process issues a voluntary page release to free a page frame that it knows it no longer needs
 - Eliminate the delay period caused by letting process gradually pass the page from its working set
 - The real hope is in compiler and operating system support




11.9 Page Size

- Some systems improve performance and memory utilization by providing multiple page sizes
 - Small page sizes
 - Reduce internal fragmentation
 - Can reduce the amount of memory required to contain a process's working set
 - More memory available to other processes
 - Large page size
 - Reduce wasted memory from table fragmentation
 - Enable each TLB entry to map larger region of memory, improving performance
 - Reduce number of I/O operations the system performs to load a process's working set into memory
 - Multiple page size
 - Possibility of external fragmentation



11.9 Page Size

Figure 11.12 Internal fragmentation in a paged and segmented system.



11.9 Page Size

Figure 11.13 Page sizes in various processor architectures.

<i>Manufacturer</i>	<i>Model</i>	<i>Page Size</i>	<i>Real address size</i>
Honeywell	Multics	1KB	36 bits
IBM	370/168	4KB	32 bits
DEC	PDP-10 and PDP-20	512 bytes	36 bits
DEC	VAX 8800	512 bytes	32 bits
Intel	80386	4KB	32 bits
Intel / AMD	Pentium 4 / Athlon XP	4KB or 4MB	32- or 36 bits
Sun	UltraSparc II	8KB, 64KB, 512KB, 4MB	44 bits
AMD	Opteron / Athlon 64	4KB, 2MB and 4MB	32, 40, or 52 bits
Intel-HP	Itanium, Itanium 2	4KB, 8KB, 16KB, 64KB, 256KB, 1MB, 4MB, 16MB, 64MB, 256MB	Between 32 and 63 bits
IBM	PowerPC 970	4KB, 128KB, 256KB, 512KB, 1MB, 2MB, 4MB, 8MB, 16MB, 32MB, 64MB, 128MB, 256MB	32 or 64 bits




11.10 Program Behavior under Paging

- Processes tend to reference a significant portion of their pages within a short time after execution begins
- They access most of their remaining pages at a slower rate



11.10 Program Behavior under Paging

Figure 11.14 Percentage of a process's pages referenced with time.



11.10 Program Behavior under Paging

- Average interfault time monotonically increases in general
 - The more page frames a process has, the longer the time between page faults



11.10 Program Behavior under Paging

Figure 11.15 Dependency of interfault time on the number of page frames allocated to a process.

